

Multi-Modal Duplicate Product Detection on Shopee

KM6312 BUBBLE TEAM

SHAN MINQIAN G2507390C

BAO CENCEN G2507353B

ZHANG JINGWEN G2507507D

QU YING G2506960G

SHEN YANBO G2503429B

TABLE OF CONTENTS

01

INTRODUCTION

02

DATA PREPROCESSING

03

TEXT MODELING

04

IMAGE MODELING

05


MULTI-MODEL
FUSION

06

EXPECTED
OUTPUTS &
BUSINESS VALUE


07

CONCLUSION



01

Introduction



Background

- Shopee's open marketplace allows multiple sellers to list the same product.
- This creates duplicate and near-duplicate listings with different titles, images, and seller info.



Key Problems

- Cluttered search results filled with repeated items.
- Fragmented reviews/ratings, weakening social proof and trust.
- Higher user effort to identify truly distinct products.
- Lower search & recommendation efficiency, hurting CTR and conversions.

Project Goal

- Build a multi-modal duplicate detection model using:
- Product titles, images, and structured attributes (price, brand, category).
- Automatically identify and group duplicate listings into unified product clusters.



Business Value

- Cleaner search experience with fewer repeated items.
- Consolidated reviews → stronger trust and higher conversions.
- Fairer exposure for sellers by reducing duplicate spam.
- Supports platform growth through improved user experience and navigation.

02

Data

Preprocessing



Dataset

- Official Shopee Product Matching (Kaggle)
- Product listings for duplicate-detection in e-commerce
- Used to train and evaluate multi-modal similarity models

Key Fields

- `posting_id` — unique listing ID
- `image` — product image file
- `title` — seller-provided title
- `label_group` — ground-truth duplicate group ID

Shopee - Price Match Guarantee

Overview Data Code Models Discussion Leaderboard Rules Team Submissions

Dataset Description

Finding near-duplicates in large datasets is an important problem for many online businesses. In Shopee's case, everyday users can upload their own images and write their own product descriptions, adding an extra layer of challenge. Your task is to identify which products have been posted repeatedly. The differences between related products may be subtle while photos of identical products may be wildly different!

As this is a code competition, only the first few rows/images of the test set are published; the remainder are only available to your notebook when it is submitted. Expect to find roughly 70,000 images in the hidden test set. The few test rows and images that are provided are intended to illustrate the hidden test set format and folder structure.

Files

`[train/test].csv` - the training set metadata. Each row contains the data for a single posting. Multiple postings might have the exact same image ID, but with different titles or vice versa.

- `posting_id` - the ID code for the posting.
- `image` - the image id/md5sum.
- `image_phash` - a perceptual hash of the image.
- `title` - the product description for the posting.
- `label_group` - ID code for all postings that map to the same product. Not provided for the test set.

Text Cleaning (title_cleaned)

- Convert to lowercase
- Remove punctuation, emojis, HTML tags
- Remove repeated spaces
- Filter invalid or empty titles
- Output: title_cleaned

Text Normalization (title_norm)

- Stopword removal
- Stemming / lemmatization
- Normalize number-unit formats (e.g., 128 GB → 128gb)
- Output: title_norm
- Used for TF-IDF, cosine similarity, and text embedding models.

```
import pandas as pd
import numpy as np
import re, random, unicodedata
import matplotlib.pyplot as plt
from pathlib import Path

RANDOM_SEED = 42
random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)

DATA_DIR = Path('.')
TRAIN_PATH = DATA_DIR / 'train.csv'
FIG_DIR = DATA_DIR / 'figures'
OUT_DIR = DATA_DIR / 'artifacts'
FIG_DIR.mkdir(exist_ok=True)
OUT_DIR.mkdir(exist_ok=True)

df = pd.read_csv(TRAIN_PATH, low_memory=False)
print('Data shape:', df.shape)
display(df.head(3))

def clean_title(t):
    t = '' if pd.isna(t) else str(t)
    t = t.lower().strip()
    t = unicodedata.normalize('NFKC', t)
    t = re.sub(r'[^\w\s\-\+]', ' ', t)
    t = re.sub(r'\s+', ' ', t)
    return t

df['title_clean'] = df['title'].apply(clean_title)
df[['title', 'title_clean']].head(5)
```

Structured Feature Construction

Brand Extraction

- Rule-based matching using a predefined brand dictionary
- Examples: vivo, Xiaomi, Samsung, adidas
- Unrecognized brands → unknown

Category Extraction

- Keyword-based mapping to categories such as
- beauty
- home
- clothes
- phone
- stationery
- Low-frequency categories merged into Others
- ⋮

⋮ New fields created: brand, category

⋮ Limitation: Shopee titles mix languages and contain many niche brands → rule-based extraction has natural constraints.

Top brands after remap:

unknown	30799
hp	1029
samsung	322
wardah	227
oppo	215
xiaomi	200
3m	154
emina	127
vivo	101
scarlett	93
apple	89
msi	68
pixy	66
loreal	61
adidas	57

Name: brand, dtype: int64

Category distribution (top 15):

others	6574
beauty	5589
home	5464
clothes	5289
phone	2212
stationery	1669
bag	1648
food	1331
shoes	974
baby	966
electronics	899
auto	861
computer	621
sports	153

Name: category, dtype: int64

Price Construction & Standardization

Shopee's raw dataset does not provide price, so we constructed it.

Simulated Price (price)

- Generated within reasonable e-commerce price ranges
- Adjusted loosely based on product type

Standardized Price (price_std)

- Z-score standardization
- Ensures consistent scaling for modeling
- New fields: price, price_std

```
Price summary (overall):
count      34250.00
mean       83.66
std        143.38
min        1.09
50%       38.98
90%       186.81
95%       327.43
max       2334.98
Name: price, dtype: float64
```

Sample rows:

	category	price	price_std
0	bag	60.030529	-0.164814
1	stationery	8.037031	-0.527452
2	food	15.687368	-0.474093
3	clothes	43.957122	-0.276921
4	food	3.101742	-0.561874
5	home	36.624555	-0.328063
6	clothes	26.993067	-0.395240
7	clothes	20.679239	-0.439277
8	home	79.197596	-0.031131
9	home	47.951989	-0.249058

Exploratory Data Analysis (EDA)

To better understand the structure and distribution of the processed dataset, we conducted three key visual analyses:

- Category distribution
- Price distribution (log scale)
- Label group size distribution

These analyses help identify imbalance, noise patterns, and structural characteristics of the dataset.




Exploratory Data Analysis (EDA)

Top 10 Product Categories

Interpretation:

- Others is the largest category due to rule-based extraction limitations & mixed-language titles.
- Major identifiable categories include beauty, home, clothes, phone, each with over 2,000–5,000 samples.
- Long-tail categories such as shoes and baby appear less frequently.
- Indicates: imbalanced category distribution, which is common in e-commerce datasets.




Exploratory Data Analysis (EDA)

Price Distribution (Log Scale)

Interpretation:

- Simulated prices follow a realistic right-skewed pattern similar to marketplace price distributions.
- Most items fall in the low-to-mid price range (1–100 units).
- A few higher prices exist but are rare, matching typical long-tail price behavior.

Using log scale reveals the distribution clearly and reduces skewness.



Exploratory Data Analysis (EDA)

Label Group Size (Duplicate Group Sizes)

Interpretation:

- Majority of label groups contain 2–3 items, meaning most duplicates are small clusters.
- A long-tail pattern exists, with a few groups containing 10–20+ items.
- Log-scale y-axis highlights the strong imbalance and sparsity of duplicate clusters.

This distribution confirms that duplicate detection is highly imbalanced, requiring careful sampling strategies.





Image Preprocessing

To ensure that all product images are consistent and usable, we applied simple but essential preprocessing steps:

- Converted all images to RGB format to unify color channels
 - Resized images to a fixed resolution (e.g., 224×224) for consistency
 - Removed unreadable or corrupted files to maintain data quality
 - Extracted basic visual features, including
 - width and height
 - average RGB values
 - grayscale brightness
- These steps make the images clean, standardized, and ready for multi-modal analysis.

03

Text Modeling



Text-only Duplicate Detection

Objective: Detect duplicate products using only textual data (title_clean)

Approach:

- Build **TF-IDF** baseline
- Apply **Sentence-BERT** for semantic matching
- Evaluate both using precision, recall, F1-score, accuracy

Model Selection & Parameter Settings

TF-IDF Baseline

Vectorizer: TfidfVectorizer(max_features=5000)

Similarity: Cosine similarity

Sentence-BERT

Model: paraphrase-multilingual-MiniLM-L12-v2

Embedding batch size: 64


Similarity: Cosine similarity via util.cos_sim()

Modeling Pipeline




Performance Comparison

Confusion Matrix - TF-IDF + pHash Baseline




Confusion Matrix - Sentence-BERT + pHash



04

Image Modeling



EfficientNet B0

EfficientNet, a landmark architecture in automated deep learning, was conceived and developed by the Google Brain team (Tan, 2020). The architecture introduces a composite scaling method that can systematically extend the benchmark model. The extension method of EfficientNet has been verified in eight models (B0-B7). The method takes the compact B0 as the **benchmark model**, and then adopts composite expansion technology to gradually build a larger-scale and stronger model sequence from B1 to B7. In mobile or edge device application scenarios with limited computing resources, EfficientNet-B0 stands out for its excellent balance of efficiency and performance. Although larger-scale variants (B1-B7) have more parameters, **B0 can show strong accuracy in tasks such as image classification and target detection, making it an ideal choice for calculating budgets** (Kansal, 2024).

This study uses a **pre-trained EfficientNetB0 model** to extract high-level visual embeddings. The Shopee product dataset contains more than **34,000 product images**. EfficientNetB0 provides a strong accuracy–efficiency trade-off and can extract embeddings for 30K+ images within reasonable runtime on a laptop environment.

Perceptual hashing (pHash)

Perceptual hashing is an advanced hash technology, which is originally designed to make the human perception system think that similar media content can produce the same or similar hash values.

Perceptual hashing's processing process include (Zauner, 2010):

- **Convert the image to grayscale first.**
- **Blurring and gamma correction**
- **Extract features and build hashes**

This project uses pHash to generate a 64-bit perceptual signature for each image, which can be used for **coarse similarity filtering** and eliminating obviously different pairs.

Image-only Experiment

EfficientNet B0

Image Preprocessing:

- Resize to 224×224 — required by CNN backbones
- Normalize pixel values to $[0, 1]$
- Convert BGR \rightarrow RGB (OpenCV default \rightarrow TensorFlow format)
- Handle corrupted and missing images

Embedding Extraction:

- Use EfficientNet-B0 without classification head
- Global Average Pooling \rightarrow Each image \rightarrow 1280-dimensional embedding vector

Output Example

- Shape: (34,250 images \times 1280 features)


Image-only Experiment

pHash

- The histogram shows an obvious multi-peak structure:
- The **vast majority** of unrelated product pairs are in the range of **22-40**.
- Small clusters of medium-similarity goods are concentrated in the range of **10-18**.
- Only a few pairing distances are **less than 10**, which strongly suggests similar duplicate or highly similar goods.

Based on this empirical distribution, we evaluate the candidate threshold from the conceptual level:


- Although the threshold below 10 can achieve extremely high accuracy, it will omit a large number of legal similar goods (low recall rate).
- A threshold higher than 18 will significantly expand the candidate pool, increase the calculation overhead and introduce excessive false alarms.
- Therefore, **selecting the threshold of Hamming distance ≤ 15** can achieve practical balance - not only retaining most visual similarity terms, but also effectively filtering irrelevant terms.



Key Findings from Image-only


To present the structure of the 1280-dimensional embedding space more intuitively, we employed t-SNE to reduce it to a two-dimensional plane. Each point corresponds to a product image and is colored according to its posting_id. Since posting_id can uniquely identify each image, the color presents a **random mixed state**, which is in line with expectations.

Although global clusters cannot be observed when coloring according to posting_id, local neighborhoods can still **capture visual similarity**, this meets the requirements of repeated commodity detection.





05

Multi-Modal Fusion



Multi-Modal Fusion

Technical Architecture Overview



Multi-Modal Fusion

Technical Architecture Overview

- End-to-end pipeline: raw CSV → deployable model
- Modular design: preprocessing, fusion, training, evaluation
- Leakage-safe split by posting_id
- Core stack: Scikit-learn, LightGBM, SBERT, EfficientNet
- Automated and reproducible workflow

```
# Step 1: Load raw data
df = pd.read_csv("shopee_dataset_final.csv")

# Step 2–3: Preprocessing + Fusion
→ SBERT + EfficientNet + Structured features → 1676–dim vector

# Step 4–5: Split + Pair generation
→ Leakage-safe split → Balanced positive/negative pairs

# Step 6–10: Training pipeline
→ CV → Tuning → Final model → Evaluation

# Step 11–14: Output & Deployment
→ Visualizations → Feature importance → Business insights → Save model
```

Python

Multi-Modal Fusion

Feature Engineering

Feature Modalities

Modality	Dimension	Implementation
Text	384	Sentence-BERT embeddings
Image	1280	EfficientNet embeddings
Structured	12	Price, brand, category, image statistics

- Unified feature vector:
1676 dimensions

Pairwise Feature Construction

```
# Core feature combinations
feature_diff = |X_i - X_j|          # Absolute difference
hadamard = X_i * X_j                # Element-wise product
cosine_sim = dot(X_i_norm, X_j_norm) # Cosine similarity
same_brand = 1 if brand_i == brand_j # Brand flag
```

Python

Multi-Modal Fusion

Model Design & Training

Leakage Prevention




Key safeguards

- 1. Split data by posting_id
- 2. Fit scaler & PCA only on training set
- 3. Strict train/test isolation

- PCA fitted only on training set

Python

LightGBM Training Workflow







Multi-Modal Fusion

Performance Analysis

- Confusion matrix shows low FP/FN
- AUC: 0.9883
- ROC curve indicates strong separability
- Accuracy: 95.44%
- Precision: 97.14%
- Recall: 93.64%
- F1-Score: 95.36%
- KDE probability distribution shows clear separation between duplicates and non-duplicates.

Model Evaluation Overview




Multi-Modal Fusion

Feature Importance Analysis

- Manhattan distance (text embeddings): strongest semantic signal for pairwise similarity
- pHash similarity (image): captures visual resemblance, useful for hard negatives
- Hadamard product (text): encodes interaction between embeddings
- Category/Brand match flags: provide strong structural cues for duplicates

Top 15 Feature Importances




Ablation Study & Insights

- Text + Numeric features yield best performance ($F1 = 0.9873$)
- Text only performs nearly as well ($F1 = 0.9866$)
- Text features dominate the overall signal
- Selective fusion is necessary

Feature Set	F1	AUC	Rank
Text + Numeric	0.9673	0.9917	1
Text Only	0.9656	0.9924	2
Text + Image	0.9565	0.9883	3
All Features	0.9560	0.9885	4
Image + Numeric	0.8798	0.9396	5
Image Only	0.8792	0.9381	6
Numeric Only	0.8718	0.9341	7

Preliminary Clustering Attempt



06

Expected Outcomes & Business Value



Commercial Problem & Motivation

- Shopee marketplace suffers from duplicate and near-duplicate product listings
- Consequences: search clutter, fragmented reviews, weak credibility, lower conversion
- Duplicate listings reduce user trust and increase bounce rate
Key metrics affected: CTR ↓, CVR ↓, GMV ↓
- Business need: automated, scalable duplicate detection → quality, fairness, revenue



Expected Outcomes

– Platform-Level Impact

AREA	OUTCOME
Search Quality	Cleaner interface, higher content relevance
User Behavior	Faster search navigation, higher trust
Engagement Metrics	CTR ↑ due to more meaningful clicks
Purchase Confidence	CVR ↑ through unified ratings & reviews
Revenue Impact	GMV ↑ due to CTR × CVR uplift

How Duplicate Detection Drives Business Metrics

Cleaner Catalog → Better Guidance → Higher Trust → More Purchases

Business Metric Impact

CTR ↑

Reduced noise → More relevant clicks

CVR ↑

Centralized reviews → Stronger purchase decision

GMV ↑

Higher CTR and CVR lead directly to revenue growth

Bounce Rate ↓

Less frustration → Higher retention


Marketplace Ecosystem & Governance Impact

- Prevents sellers from posting duplicate listings to manipulate exposure
- Ensures fair visibility for authentic and small sellers
- Strengthens platform governance and policy enforcement
- Reduces customer disputes and misleading product claims
- Builds a trusted and sustainable marketplace ecosystem

Operational and Strategic Benefits

Area	Benefit
Moderation Cost	Automated duplicate detection reduces manual review burden
Efficiency	Approves quality listings faster, filters fake/repeated uploads
Content Quality	Standardizes catalog, improves metadata consistency
Scalability	Suitable for real-time deployment at platform level

Business Value Chain



- Shows how technical model links to business metrics & ecosystem health
- Demonstrates cascading value from search quality to marketplace trust
- Highlights both revenue and governance dimensions

07

Conclusion




Final Takeaway: What This Project Proves

Multimodal fusion (Text + Image + Metadata) outperforms single-source detection

- pHash improves candidate retrieval speed and complements deep models
- Duplicate detection is not only a technical challenge, but a commercial enabler
- Direct impact on: search quality, CTR, CVR, GMV, fairness, and governance
- Supports real-world deployment in Shopee-like e-commerce environments

Future Deployment Potential





THANKS

